

# Security and Compliance Review of HealthPrep V2

## TLS and Data-in-Transit Security

Secure TLS encryption is critical for protecting data in transit. The application is already configured to use HTTPS-only cookies and safe defaults: - **Secure Cookies:** Session cookies are marked Secure and HttpOnly, and SameSite is set (Lax) <sup>1</sup> to prevent theft or XSS leaks. CSRF protection is enabled by Flask-WTF <sup>2</sup>. - **Proxy Awareness:** The app applies Flask's `ProxyFix` middleware <sup>3</sup> so it correctly detects the original protocol when behind a proxy (like Replit's platform) - this ensures URLs and redirects remain on `https://`.

**Improvement Opportunities:** - **Enforce HTTPS:** Implement automatic redirects from HTTP to HTTPS in production (e.g. using HSTS headers or Flask-Talisman) to ensure all traffic is encrypted. Currently, no explicit redirect to HTTPS is in place (relying on Replit's default). - **Strong TLS Configuration:** If hosting outside Replit, configure the web server for TLS 1.2+ (disabling older protocols) and enable HSTS (Strict-Transport-Security) so browsers only connect via HTTPS going forward. - **Certificate Handling:** Ensure that any external API requests (e.g. to Epic systems) use HTTPS endpoints (which they do by default <sup>4</sup>) and verify certificates. Do not disable SSL verification in HTTP clients. This prevents man-in-the-middle attacks on external data flows.

Overall, the application is well-positioned for TLS usage (thanks to secure cookies and proxy handling), but explicit enforcement and hardened TLS settings will further strengthen in-transit security.

## HIPAA Compliance and BAA Considerations

HealthPrep V2 includes features aimed at HIPAA compliance. It has a dedicated security configuration for HIPAA <sup>5</sup> and robust auditing of sensitive data access. Notable measures in place: - **Audit Logging:** All admin actions and any PHI access are logged with user, timestamp, and context. The `AdminLog` model tracks the event type, user, org, IP, and even patient or document IDs accessed <sup>6</sup>. For example, the system logs whenever patient data is viewed or a document is downloaded <sup>7</sup>, creating an audit trail for HIPAA. - **Retention Policies:** Log retention is configured for 7+ years by default <sup>8</sup> <sup>9</sup>, meeting HIPAA's minimum retention recommendation. Each organization record even stores an `audit_retention_days` (default ~2555 days ≈ 7 years) <sup>9</sup>. - **PHI Filtering:** The code defines regex patterns to identify PHI like SSNs, phone numbers, etc. <sup>10</sup>. This suggests the system can redact or flag PHI in outputs. (Ensure these filters are applied wherever logs or non-authorized views might contain PHI.)

**Improvement Opportunities:** - **Hosting Environment (BAA):** Replit is not a HIPAA-compliant hosting environment and likely will not sign a Business Associate Agreement. For production use with real PHI, migrating to a HIPAA-compliant host (AWS, Azure, GCP or a specialized provider) is critical. The new host should sign a BAA and provide encryption at rest, network isolation, and monitoring. - **Encryption at Rest:** Consider encrypting sensitive data in the database. Currently, PHI like patient names, etc., are stored in plaintext. While a secure DB and host might suffice, encrypting fields (or the entire database volume) adds a

layer of protection if the storage is compromised. In particular, the Epic client secrets and tokens should be stored encrypted (more on this below). - **Audit Log Integrity:** Ensure audit logs are tamper-proof. The application logs to the database <sup>11</sup>; restrict who can modify these records. You might implement append-only logging or at least database-level permissions so that even admins cannot alter past logs. Regularly back up and archive logs in a secure, immutable form to meet compliance. - **Policy and Procedure Compliance:** Beyond code, maintain HIPAA policies: role-based access (addressed below), training for users, breach notification processes, etc. Technically, features like automatic log-off (the system uses a 30-minute inactivity timeout by default) and account lockout are in place – ensure these align with your org's policy (you might adjust timeout duration or lockout thresholds as needed). - **Regular Audits:** Periodically conduct security audits or penetration tests on the system. Verify that no PHI is unintentionally leaking (e.g., in logs or error emails). The PHI filtering can be used to scan logs for any unredacted PHI and verify compliance.

In summary, the software includes several HIPAA-oriented safeguards (logging, retention, filtering) <sup>12</sup> <sup>6</sup>. The primary improvement is to ensure the hosting environment and operational practices (BAA, encryption, monitoring) are brought up to compliance standards, since the code features alone are not sufficient for full HIPAA compliance.

## Database Storage and Multi-Tenancy Security

The application employs a multi-tenant database design, meaning data for different client organizations is isolated by an `org_id` field: - **Organization Data Isolation:** Most models (Patient, Document, Screening, etc.) include an `org_id` foreign key. The User model has a unique constraint on (username, org\_id) and (email, org\_id) <sup>13</sup> to prevent cross-tenant duplicates. This ensures two orgs could have a user with the same username without conflict. - **Query Filtering:** The code uses utilities to enforce org-based filtering. For example, the `require_organization_access` decorator checks that the current user's org matches the org in question, aborting if not <sup>14</sup>. In practice, this means an authenticated user cannot craft a URL to access another organization's data – the backend will return 403 Forbidden <sup>14</sup>. Additionally, many queries explicitly filter by `current_user.org_id`; e.g., admin dashboard queries only fetch logs and documents for the user's own org <sup>15</sup> <sup>16</sup>. - **Role Separation for Tenancy:** Regular admins (non-root) are always tied to an `org_id`, and they only manage users and data within that org. The root admin (superuser) has `org_id = None` and is handled separately (see below). This design maintains tenant data segregation effectively.

**Improvement Opportunities:** - **Enforce Org Scope Everywhere:** Double-check that *all* data queries are wrapped in an org filter. The utilities exist, but it's worth auditing controllers for any missing filters. For instance, older code (like `admin_routes_backup.py`) did not filter logs by org <sup>17</sup> – the new code does <sup>18</sup>. Ensure no endpoints use `User.query.all()` or similar without `org_id` criteria (the current code seems to handle this well). - **Database-Level Isolation:** For an extra layer, you could implement database row-level security or separate schemas/databases per tenant. Given the current size, the single database with `org_id` is fine, but as you scale, separate schemas or even separate DB instances per client might be considered for stricter isolation and easier per-tenant data dumps. - **Data Deletion and Retention per Org:** Provide mechanisms to cleanly delete or export an organization's data if needed. The code provides a route to delete an organization <sup>19</sup>, but it first requires removing users <sup>20</sup>. When an org is deleted, thanks to `cascade='all, delete-orphan'` on relationships, related data (patients, documents, etc.) should be removed <sup>21</sup>. Test this process to ensure no orphaned PHI remains after an org deletion. Also consider a “soft delete” or backup when deleting orgs, in case it's done accidentally. - **Database Credentials & Access:** Ensure the database itself is secured. If using PostgreSQL, use a strong password for the DB user and

restrict network access (if the DB is cloud-hosted, use firewall / security groups to allow only the app server). On Replit, if it's SQLite, ensure the `.db` file isn't exposed via web routes. The file is stored under an `instance/` folder by default <sup>22</sup>, which Flask doesn't serve, so that's good. - **Backups:** Implement regular backups of the database. This can be as simple as a cron job to dump the DB to an encrypted storage location. Given HIPAA, you'd want to be able to restore data in case of any corruption or ransomware scenario.

The multi-tenancy strategy in HealthPrep V2 is solid – it clearly separates tenant data and includes checks to prevent unauthorized cross-access <sup>14</sup> <sup>13</sup>. The focus should be on maintaining those checks as code evolves, and on securing the database at an infrastructure level.

## User Roles and Access Control

The system defines four main user roles: **Root Administrator**, **Organization Admin**, **Medical Assistant (MA)**, and **Nurse**. The authorization logic and UI flows ensure that each role's access is limited appropriately: - **Root Admin (Superuser):** This role (`is_root_admin=True`) can manage all organizations and system-wide settings. Root admins have their own dashboard (at `/root-admin/*`). The code uses a decorator to restrict root-admin pages – e.g., `@root_admin_required` checks `current_user.is_root_admin_user()` before allowing access <sup>23</sup>. Similarly, a normal admin is redirected away from root-admin URLs and vice versa, by a safety check on each request <sup>24</sup>. This prevents, say, an org admin from even trying to load a root-admin page. - **Org Admin:** Can manage everything within their organization (users, Epic credentials, patient data). The `@admin_required` decorator is used on admin routes to ensure the user is an admin of some sort <sup>25</sup>. Internally, `is_admin_user()` returns true for either role 'admin' or a root admin <sup>26</sup>, so root can access admin routes too. Notably, org admins cannot affect other organizations – all such routes enforce org scope (as discussed above). - **MA and Nurse:** These are more limited roles (presumably can view and update patient screenings but not manage users or org settings). The UI and templates likely hide admin functions from them. The backend also checks `current_user.is_admin_user()` for admin endpoints, which a nurse/MA would fail, thus getting a redirect or 403. Additionally, the `require_role('admin')` type decorators could be used to enforce specific roles on certain actions <sup>27</sup> (ensuring only admins do X, or only MAs do Y if implemented).

**Security Measures in Place:** - **Dashboard Segregation:** The app's `before_request` logic redirects users to the appropriate dashboard if they try to access an area outside their role <sup>24</sup>. For instance, if a nurse somehow navigates to an admin URL, they'll be sent to the user dashboard. - **Immunity of Superuser Accounts:** The UI prevents critical actions on root admin accounts. For example, the Root Admin user cannot be deactivated via the toggle in the interface <sup>28</sup>, nor deleted. Likewise, the last admin in an organization is protected from deletion in the UI (a warning is shown) <sup>29</sup>. These guardrails ensure you don't accidentally lock yourself out of an org or the system. - **Audit of Privileged Actions:** All user management actions are logged. Creation, editing, deletion of users and orgs call `log_admin_event` with details <sup>30</sup> <sup>31</sup>. This provides accountability – if a root admin changes a role or an admin deletes a user, there's an audit trail.

**Improvement Opportunities:** - **Two-Person Integrity for Destructive Actions:** For very sensitive operations (e.g., deleting an entire organization or mass patient deletion), you might require a secondary confirmation or approval. Currently, a single admin or root admin can perform these actions. Implementing an optional “Are you sure? Type the org name to confirm” or requiring another admin's approval could

prevent mistakes or malicious purges. - **Finer-Grained Permissions:** As the system grows, consider a more granular RBAC. For example, within the “admin” role, you might later differentiate between a clinical admin who manages screenings vs. an IT admin who manages users. Right now roles are coarse. The code is already structured to check for roles on endpoints <sup>27</sup>, so extending roles or permissions (perhaps via a permissions table) in the future is feasible. - **User Impersonation Logging:** If root admins can impersonate org users (not sure if this feature exists – it’s not explicitly in the code), that should be logged. If it doesn’t exist, ignore this point. If it ever is added as a support tool, make sure to track “Root Admin X acting as User Y”. - **UI Feedback on Permissions:** Ensure that menu items and buttons not usable by a nurse or MA are hidden or disabled. This seems to be done (different base templates for user vs admin interfaces), but double-check there are no broken links where a nurse could attempt an admin action (they’d be stopped server-side, but better to not tempt them at all). - **Prevent Orphaned Orgs/Users:** The system should ensure that every organization always has at least one active admin. The UI warns about not deleting the last admin, but the `delete_user` route on the backend does not explicitly prevent it <sup>32</sup>. It would be wise to add a server-side check: if `user.role == 'admin'` and they are the only admin in that org, refuse deletion or deactivation. This ensures every org retains an admin to manage it. (Root can always step in via root-admin panel, but it’s cleaner if orgs are self-contained.)

Overall, role-based access is well-implemented – critical admin functions are limited to the right roles and clearly separated in both code and interface <sup>24</sup> <sup>25</sup>. Enhancements should focus on edge cases (last admin removal, etc.) and future-proofing for more complex permission scenarios.

## Password Policies and Authentication

The application takes password security seriously and follows best practices in many areas: - **Strong Password Requirements:** By default, passwords must be at least 8 characters and include an uppercase letter, lowercase letter, number, and special character <sup>33</sup>. The `validate_password` function enforces this and returns errors if criteria are not met <sup>34</sup>. This helps encourage strong passwords for all users. - **Secure Password Storage:** Passwords are never stored in plaintext. The User model hashes passwords with Werkzeug’s PBKDF2 algorithm (using `generate_password_hash`) <sup>35</sup>. When checking login, it uses `check_password_hash` to verify credentials <sup>36</sup>. This means even if the database were compromised, attackers get only hashed passwords (which are computationally expensive to crack, especially with PBKDF2 salting and stretching). - **Account Lockout Mechanism:** The system defends against brute-force attacks by locking accounts after 5 failed login attempts. Specifically, `User.record_login_attempt()` will set a `locked_until` timestamp 30 minutes in the future after 5 failures <sup>37</sup>. Any login attempts while the account is locked are rejected. This is a solid safeguard against password guessing. - **Session Timeout:** By default, sessions expire after 30 minutes of inactivity (each user has `session_timeout_minutes=30`) <sup>38</sup>. The app updates `last_activity` on each request <sup>39</sup> and logs out users who have been idle too long. This reduces the risk of someone walking away from a workstation and an attacker hijacking their session. - **Additional Safeguards:** The app sets the Flask-Login `SESSION_PROTECTION` (not explicitly seen, but likely default) and already has CSRF protection enabled for logins and forms. Also, the `is_active_user` flag on accounts allows easy deactivation of accounts if a user should no longer log in <sup>40</sup> <sup>41</sup>.

**Improvement Opportunities:** - **Two-Factor Authentication:** Enable the `two_factor_enabled` field that exists on the User model <sup>42</sup> by implementing 2FA flows. For admins or any user dealing with PHI, a TOTP (authenticator app code) or SMS-based second factor would significantly improve security. This could be

optional per user but highly recommended for root and org admins. The groundwork (field in DB) is there, but you'd need to integrate a library or API for sending codes and verifying them at login. - **Password Hash Algorithm:** Werkzeug's default PBKDF2-HMAC-SHA256 is secure, but consider using an even stronger algorithm like bcrypt or Argon2 for new passwords, especially if supported by Flask security plugins. PBKDF2 is fine and NIST-approved; this is a mild suggestion for future-proofing. If changing, you'd need a migration path (e.g., on login, detect old hash format and re-hash). - **Password Rotation Policy:** Currently, there's a flag in the admin settings UI for "require password change" <sup>43</sup>. If you plan to enforce periodic password changes (some organizations require e.g. every 90 days), you could implement tracking of last password change and prompting users. Note however that forced rotation is debated in security – it can sometimes lead to weaker passwords – so use only if policy demands. - **Login Notification:** As an extra layer, consider notifying users (via email) on successful logins or on account lockouts. For instance, "Your account was locked due to 5 failed attempts" or "New login from IP X". This can help detect unauthorized access attempts early. - **Protect Admin Accounts:** Ensure admin/root usernames are not easily guessable. Currently the default root username is "rootadmin" (or was, during setup) and an example admin is "admin". In production, encourage using less generic usernames for these accounts, so attackers can't easily target them. Also, if not already done, rename or remove any default credentials (see below) so there are no known passwords in the system.

On the whole, authentication in the app is robust – strong password rules <sup>33</sup>, hashed storage <sup>35</sup>, lockouts <sup>44</sup>, and timeouts are all in place. By adding 2FA and continuing to monitor login attempts (the app does log them), you can significantly harden account security against modern threats.

## Credential Management and Secrets Handling

Proper handling of secrets (API keys, passwords, tokens) is crucial. Currently, there are both good practices and areas to improve in HealthPrep V2:

**Current State:** - **Use of Environment Variables:** The app expects critical secrets to be provided via environment. For example, the Flask `SECRET_KEY` is taken from an env var if set <sup>45</sup>, and the database URL is from `DATABASE_URL` env var <sup>46</sup>. The Epic OAuth client ID/secret can be set per organization through the UI rather than being hard-coded. Additionally, the JWT private keys for Epic backend auth are read from Replit Secrets (environment vars) rather than stored in code or repo <sup>47</sup>. These are all good practices – no plaintext credentials are hard-coded in the repository (except some defaults for development). - **App Secret Key:** In development, a default secret (`'dev-secret-key-change-in-production'`) is provided <sup>45</sup>, but the expectation is to override this in production. This key secures session cookies and JWT signing. Using an env var for it in production is correct; one just must be sure to actually set a strong random value. - **Epic OAuth Credentials:** Each Organization has fields for `epic_client_id` and `epic_client_secret` in the database <sup>48</sup>. This allows storing unique credentials for each client org's Epic integration. However, currently the client secret is stored in plaintext in the DB. When an org admin enters it on the Epic registration page, it's saved directly <sup>49</sup>. The code comments even note "Should be encrypted in production" <sup>50</sup>. - **Token Storage:** OAuth access tokens and refresh tokens from Epic are stored in the `EpicCredentials` table (fields `access_token`, `refresh_token`) <sup>51</sup>. These too are stored as plaintext (albeit likely short-lived tokens). - **Key Rotation Mechanism:** A positive: for the JWT client auth, the system supports key rotation by design. It can hold multiple RSA keys (with naming like `NP_KEY_2025_08_A`) and will pick the first available for signing <sup>52</sup>. The JWKS service publishes all public keys <sup>53</sup>. This means you can periodically add a new key via env, and

remove the old one, to rotate keys without downtime – a great security feature. The instructions in `utils/jwt_utils.py` explicitly guide how to add new keys and retire old ones <sup>54</sup> .

**Improvement Opportunities:**

- **Encrypt Sensitive Fields in DB:** Implement encryption for `Organization.epic_client_secret` (and possibly the tokens in `EpicCredentials`). For example, you could use a symmetric cipher (like Fernet from the cryptography library) with an app-specific key to encrypt these values before saving. The decryption key could be derived from the master `SECRET_KEY` or a separate env var. This way, if the database is compromised, the attacker cannot directly use the stolen client secrets or tokens. Given the comment in code <sup>50</sup> , this is already a known to-do.
- **Do Not Display Secrets in UI:** Currently, the Epic client secret is rendered back to the admin in the form (masked, but present) <sup>55</sup> . A savvy user could inspect the HTML or use browser dev tools to reveal the actual value. It's safer to not send the secret back to the frontend at all. Instead, you can show it as blank or obfuscated ("\*") **and only update it server-side if a new value is provided. This prevents any possibility of the secret leaking via the web UI.**
- **\*Ensure Default Secrets Are Replaced:** Double-check that in the live environment no default secrets are in use. The `SECRET_KEY` must be a secure random value (not the "dev-secret-key"). The code's default admin credentials (username "admin" with password "admin123" for example) <sup>56</sup> should not exist in production data – if they do, remove or change them immediately. Similarly, the sample root admin created by the setup script (username "rootadmin" / "rootadmin123") <sup>57</sup> should be changed. In short, any credential that appears in documentation or code should be considered compromised – generate new, unique ones for actual use.
- **Secret Storage Management:** Continue using environment variables (or a secure secrets manager) for all keys. In Replit, the Secrets feature is being used <sup>47</sup> – that's good as it doesn't expose values in code. If migrating to another host, use an equivalent mechanism (AWS SSM Parameter Store, Docker secrets, etc.). Avoid ever logging these secrets. Scour the logs to ensure no accidental `print()` of a secret occurred (the code logs a lot of info, but I didn't see secrets being logged except maybe Epic token exchange failure might log an error containing part of a URL; ensure no secrets in query params).
- **Session Cookie Secret Rotation:** Flask's session cookies are signed with `SECRET_KEY`. Rotating that key will invalidate all sessions (which is acceptable if planned). Consider rotating this key if you suspect it's compromised, with communication to users that they will need to log in again. There's no automated rotation in Flask for this, but you can manually change it periodically (e.g., annually) as a hygiene practice.
- **API Credentials Rotation:** For the Epic client ID/secret, have a procedure for rotating those if needed (Epic might issue new secrets if old ones are compromised). The system allows updating them via the UI. Ensure that if a new secret is saved, the old tokens are invalidated (possibly Epic will do that on their end). You might want to proactively delete any stored access/refresh tokens whenever the client credentials change, forcing a fresh OAuth flow with the new credentials.

By encrypting at-rest secrets and tightening how they are handled in transit, we can significantly reduce the risk of credential leakage. The app already uses environment-driven secrets and has support for key rotation in some areas <sup>54</sup> <sup>58</sup> ; extending these practices to all sensitive info will strengthen overall security.

## Codebase and Redundancy Cleanup

Maintaining a clean codebase helps avoid security bugs and confusion. A few housekeeping improvements can be made:

- **Remove Deprecated Files:** The repository contains old/backup files, such as `admin_routes_backup.py` <sup>59</sup> . This duplicate code (from before multi-tenancy updates) could confuse developers or, worst-case, be accidentally used if someone imports it. It's best to remove or clearly mark such files as deprecated. Rely on the updated `admin_routes.py` going forward.
- **Consolidate Migration Scripts:** There are one-off scripts like `migrate_root_admin.py` and `migrate_multi_tenancy.py`.

After running these once (to alter the schema for root admin and multi-tenancy), consider integrating their changes into your primary migration system (Flask-Migrate/Alembic). If the database is already migrated, you might remove these scripts or move them to a `/migrations/archive` folder. This prevents someone from running them again by mistake. At minimum, ensure they are idempotent. (For example, `migrate_root_admin.py` checks if the column exists before adding <sup>60</sup> – that's good.) - **Default Data and Accounts:** As noted, the code seeds some default users (admin/user with simple passwords) for development convenience <sup>56</sup>. Make sure these lines do not execute in a production context. It appears they are part of a function `initialize_default_data()` that isn't called unless manually invoked. It would be wise to remove or disable such dev helpers in the production branch, so nobody accidentally runs them. The same goes for the root admin setup that auto-creates a `rootadmin/rootadmin123` user <sup>57</sup> in an interactive script – after the real root admin is created, delete or secure that script. - **Config Management:** Verify that the app is running with the correct config (ProductionConfig vs DevelopmentConfig). In `config/settings.py`, ProductionConfig sets `DEBUG=False` which is important to prevent debug info leak <sup>61</sup>. Ensure the environment variable (e.g., `FLASK_ENV`) is set to production in deployment. Also, consider using `SESSION_COOKIE_SECURE=True` even in development if possible, to closely mimic production (currently dev sets it False <sup>62</sup>, which is okay). - **Resource Files:** Remove any test data or placeholder files that aren't needed. For example, if there are sample PDFs or images used during dev that contain fake patient data, eliminate them from the live codebase. - **Logging Noise:** The application logs a lot of info (which is great for debugging). For production, consider adjusting log levels or removing overly verbose logging once stable. For instance, logging full HTTP headers or large JSON payloads could inadvertently log PHI. Make sure the logging configuration (currently basicConfig at INFO level <sup>63</sup>) is appropriate. You might switch to WARNING level in production and direct logs to a secure location with log rotation in place.

By cleaning up the codebase and configs, you reduce the attack surface (less dead code that could have vulnerabilities) and ensure that only intended functionality is running. It also simplifies compliance audits, since there are fewer “mystery scripts” lying around.

## Epic OAuth Integration and Organizational Data

Integrating with Epic's FHIR API requires careful handling of credentials and organizational context. Current implementation aspects and improvements are as follows:

**Current Implementation:**

- **Per-Organization Credentials:** The system supports storing distinct Epic OAuth credentials for each organization. The Organization model holds the client ID, client secret, FHIR base URL, and other necessary info like OAuth endpoints and organization identifiers <sup>64</sup> <sup>65</sup>. This means each tenant can connect their own Epic instance or sandbox. The Epic registration admin UI collects these fields from an org admin <sup>66</sup> <sup>55</sup>.
- **SMART on FHIR Compliance:** The integration follows SMART on FHIR specs – using `authorization_code` grant by default. The `FHIRClient` class constructs the auth and token URLs based on the provided FHIR base URL <sup>67</sup> and requests the proper scopes (including `offline_access` for refresh tokens) <sup>68</sup>. The code handles exchanging auth code for token <sup>69</sup> <sup>70</sup> and refreshing tokens when expired <sup>71</sup> <sup>72</sup>. All HTTP calls to Epic are made over HTTPS using the `requests` library (which verifies SSL by default).
- **JWT Client Credentials (System Integration):** The app is prepared for backend services authentication to Epic (for scenarios without user interaction). It uses JWT assertion (with private keys, JWKS, etc.) as discussed. There's a health check for key consistency <sup>73</sup> and a `.well-known/jwks.json` endpoint to allow Epic to fetch public keys. This indicates a thorough implementation for Epic's system-level OAuth variant as well.
- **Organization Context:** When making API calls, the code always knows which

organization it's dealing with (e.g., `EpicFHIRService(organization_id)` ensures it uses the right client credentials <sup>74</sup>). The app tracks the connection status per org (`Organization.is_epic_connected`, `last_epic_sync`, etc.) <sup>75</sup> and even logs errors per org to help with troubleshooting. This granular status tracking is helpful to ensure one tenant's Epic issues do not affect others. - **Epic OAuth User Creation:** The question mentioned user creation on `/admin/epic/registration` by an org admin. Indeed, an org admin can create the credentials there, and the required fields are username, password, email for creating an Epic account? (It sounds like they might create an Epic-linked user account in our system.) In the current code, user creation in general requires those fields <sup>76</sup>, and root admins can assign roles when creating users <sup>77</sup> <sup>78</sup>. All roles share the same secure practices described earlier, so there's no special vulnerability for Epic-linked users beyond normal user account security.

**Improvement Opportunities:** - **Encrypt and Secure Epic Credentials:** As noted, encrypting the `epic_client_secret` in the database is important. Also consider encrypting access/refresh tokens in `EpicCredentials` table, since these tokens can grant access to PHI on Epic. At minimum, treat the refresh token like a password – it should be stored hashed or encrypted if possible (though hashing a token isn't practical if you need to use it; encryption with a key stored server-side is the way). - **Least Privilege for Epic API User:** If the Epic integration uses a service account or specific user account on the Epic side, ensure that account has minimal permissions (only the FHIR scopes needed, which looks to be the case <sup>68</sup>). Do not request more data than necessary. The scopes listed (Patient/Observation/Condition, etc.) seem appropriate for a screening app. - **Epic Environment Separation:** The code distinguishes sandbox vs production environments for Epic (e.g., separate client IDs, different base URLs, and `epic_environment` flag) <sup>79</sup>. Make sure that in production, `epic_environment` is set to "production" for live orgs, and that the corresponding URLs (auth and token endpoints) are correctly stored. The system default is geared to Epic's public sandbox <sup>80</sup>. For real clients, you'll use their unique endpoints (which the model has fields for <sup>65</sup>). Test the flow in a non-prod and prod setting to ensure the dynamic URL substitution works. - **Monitoring and Error Handling:** Given that Epic integrations can fail for various reasons (expired credentials, downtime, etc.), ensure robust monitoring. The app does log errors to `Organization.last_epic_error` <sup>81</sup>. It might be useful to surface those errors to the org admin in the UI (perhaps it already does via `epic_connection_status`). Also consider email alerts or a dashboard for the root admin to see if any org's Epic sync has been failing (so you can proactively reach out). - **Epic Client Secret Changes:** If an organization needs to update their Epic client secret (say it was rotated or exposed), they can do so in the UI. After updating, ensure the application clears or re-establishes any existing sessions/tokens with Epic. The code currently just updates the secret and leaves tokens as is <sup>82</sup>. If the old secret is invalid, token refresh will start failing – the org admin might need to re-run the OAuth authorization. It may be wise to detect changes to client ID/secret and force a fresh authorization (perhaps by marking `is_epic_connected=False` to prompt re-connection). - **Data Mapping and Validation:** Make sure the system properly uses the `epic_endpoint_id`, `epic_organization_id` fields if required for Epic's App Orchard registration. These might need to be sent or used in API calls (depending on Epic's requirements). If they are not needed, ensure they're at least captured for reference. The code currently doesn't show usage of these fields, so verify if Epic requires them for any calls or setup. - **Separation of Concerns:** The Epic integration code is quite complex and is heavily integrated into the app. In the future, consider isolating it in a module or service with clear interfaces. This is more of a maintainability note: it's already partly done with `services/epic_fhir_service.py` and `emr/fhir_client.py`, which is good.



All in all, the Epic OAuth integration is well thought-out in this software – it handles both interactive and service workflows, per-tenant credentials, and has monitoring points. The main security focus should be on protecting the OAuth secrets/tokens and ensuring that a breach in one tenant's Epic credentials cannot compromise others (which, thanks to multi-tenancy isolation, is largely achieved). Regular reviews of Epic audit logs (both on our side and Epic's side) are recommended once live, to verify that access is only happening as expected.

## Deployment and Replit Resource Management

Running the application on Replit has been convenient for development and testing, but there are some considerations for security, compliance, and performance:

- **HIPAA on Replit:** As mentioned, Replit is not certified for HIPAA compliance. If you continue to use Replit for demos or development, avoid real PHI in that environment. For production, migrating to a compliant host is advised. If you did want to stick with Replit (not recommended for prod PHI), you would need a signed BAA which Replit likely won't provide, and you'd have to ensure data encryption and network isolation manually – not trivial on a multi-tenant platform.
- **Data Persistence:** Replit projects have an ephemeral nature – though they do save file state, the instance can shut down when idle (unless on a paid always-on plan). The team added a `keep_alive.py` script<sup>83</sup><sup>84</sup> to ping the app periodically and keep it running during critical processes (like Epic's app approval which might involve external server checks to your `.well-known/jwks.json`). This is a clever workaround, but for production you'd want a more robust hosting where uptime is guaranteed and such ping hacks aren't needed.
- **Storage Limits:** Replit imposes limits on storage and memory. Monitor the size of the SQLite database (`healthprep.db`) and any uploaded documents. The config limits file uploads to 16MB each<sup>85</sup>, which is good, but many large files could still accumulate. Consider periodically cleaning out or archiving old files. If the app will store a lot of documents, moving to a dedicated storage (S3 or similar) would be more scalable.
- **Performance:** Replit containers have limited CPU/RAM. If you onboard many organizations and process a lot of data (OCR, etc.), you may hit performance issues. The code uses background tasks (RQ) for heavy lifting, but on Replit you may not have a persistent worker process for RQ. Watch memory usage especially during OCR or batch processing.
- **Environment Secrets:** Continue using Replit's Secret Manager for all config (DB URL, email credentials if any, etc.). The JWT keys are stored there as shown<sup>47</sup>. This keeps them out of source control. If moving to another host, use an analogous mechanism (don't hardcode secrets into docker images or code).
- **Private vs Public Repl:** Ensure your Repl is private if it contains any sensitive info. A public Repl's code (and potentially environment variables) might be visible to others. Only team members should have access. If using version control like GitHub, never push the `.env` file or any secret values.
- **Resource Monitoring:** Implement logging or alerts for when the app approaches resource limits. For example, if using SQLite, log the database file size periodically. Also log memory usage if possible. It's better to catch a looming outage (due to out-of-memory or disk-full) before it happens.
- **Backup Strategy on Replit:** Replit doesn't automatically backup your database. You could create a simple backup script to periodically copy `healthprep.db` out to a secure storage (perhaps commit to a private GitHub repo or upload to cloud storage, encrypted). This ensures you don't lose data if the Repl is corrupted or you make a mistaken deployment.

In summary, Replit is a great development sandbox but not intended for secure, high-reliability production use – particularly not for healthcare data. Plan for a transition to a more robust environment as the project moves forward. In the meantime, use Replit’s features (Secrets, always-on, etc.) to mitigate risks and keep a close eye on resource usage and data persistence.

---

By addressing the points above – enforcing TLS, shoring up HIPAA compliance, securing stored credentials, tightening multitenancy, and cleaning up any weak links – HealthPrep V2 will be much better positioned in terms of security and compliance. It’s clear that the application was built with these concerns in mind (strong password policy <sup>33</sup>, audit logs, secure defaults, etc.), so the recommendations mainly involve extending or refining those existing practices. Prioritize the high-impact items like environment/BAA and encryption of sensitive data. With those improvements, you’ll significantly “immunize” the system against data breaches and unauthorized access while maintaining the functionality of the screening engine.

---

1 2 5 8 10 12 33 34 **security.py**

<https://github.com/mitfusco98/HealthPrepV2/blob/ec308aa61c685848a68ae5060ad34ca858be31c/config/security.py>

3 22 24 63 **app.py**

<https://github.com/mitfusco98/HealthPrepV2/blob/ec308aa61c685848a68ae5060ad34ca858be31c/app.py>

4 67 69 70 **fhir\_client.py**

[https://github.com/mitfusco98/HealthPrepV2/blob/ec308aa61c685848a68ae5060ad34ca858be31c/emr/fhir\\_client.py](https://github.com/mitfusco98/HealthPrepV2/blob/ec308aa61c685848a68ae5060ad34ca858be31c/emr/fhir_client.py)

6 9 11 13 21 26 35 36 37 38 39 42 44 48 51 64 65 75 79 81 **models.py**

<https://github.com/mitfusco98/HealthPrepV2/blob/ec308aa61c685848a68ae5060ad34ca858be31c/models.py>

7 14 27 50 **multi\_tenancy.py**

[https://github.com/mitfusco98/HealthPrepV2/blob/ec308aa61c685848a68ae5060ad34ca858be31c/utils/multi\\_tenancy.py](https://github.com/mitfusco98/HealthPrepV2/blob/ec308aa61c685848a68ae5060ad34ca858be31c/utils/multi_tenancy.py)

15 16 18 25 30 31 32 40 41 43 76 77 78 **admin\_routes.py**

[https://github.com/mitfusco98/HealthPrepV2/blob/ec308aa61c685848a68ae5060ad34ca858be31c/routes/admin\\_routes.py](https://github.com/mitfusco98/HealthPrepV2/blob/ec308aa61c685848a68ae5060ad34ca858be31c/routes/admin_routes.py)

17 59 **admin\_routes\_backup.py**

[https://github.com/mitfusco98/HealthPrepV2/blob/ec308aa61c685848a68ae5060ad34ca858be31c/routes/admin\\_routes\\_backup.py](https://github.com/mitfusco98/HealthPrepV2/blob/ec308aa61c685848a68ae5060ad34ca858be31c/routes/admin_routes_backup.py)

19 20 23 82 **root\_admin\_routes.py**

[https://github.com/mitfusco98/HealthPrepV2/blob/ec308aa61c685848a68ae5060ad34ca858be31c/routes/root\\_admin\\_routes.py](https://github.com/mitfusco98/HealthPrepV2/blob/ec308aa61c685848a68ae5060ad34ca858be31c/routes/root_admin_routes.py)

28 29 **edit\_user.html**

[https://github.com/mitfusco98/HealthPrepV2/blob/ec308aa61c685848a68ae5060ad34ca858be31c/templates/root\\_admin/edit\\_user.html](https://github.com/mitfusco98/HealthPrepV2/blob/ec308aa61c685848a68ae5060ad34ca858be31c/templates/root_admin/edit_user.html)

45 46 56 61 62 80 85 **settings.py**

<https://github.com/mitfusco98/HealthPrepV2/blob/ec308aa61c685848a68ae5060ad34ca858be31c/config/settings.py>

47 54 **jwt\_utils.py**

[https://github.com/mitfusco98/HealthPrepV2/blob/ec308aa61c685848a68ae5060ad34ca858be31c/utils/jwt\\_utils.py](https://github.com/mitfusco98/HealthPrepV2/blob/ec308aa61c685848a68ae5060ad34ca858be31c/utils/jwt_utils.py)

49 68 **epic\_registration\_routes.py**

[https://github.com/mitfusco98/HealthPrepV2/blob/ec308aa61c685848a68ae5060ad34ca858be31c/routes/epic\\_registration\\_routes.py](https://github.com/mitfusco98/HealthPrepV2/blob/ec308aa61c685848a68ae5060ad34ca858be31c/routes/epic_registration_routes.py)

52 58 73 **jwt\_client\_auth.py**

[https://github.com/mitfusco98/HealthPrepV2/blob/ecb308aa61c685848a68ae5060ad34ca858be31c/services/jwt\\_client\\_auth.py](https://github.com/mitfusco98/HealthPrepV2/blob/ecb308aa61c685848a68ae5060ad34ca858be31c/services/jwt_client_auth.py)

53 **jwt\_service.py**

[https://github.com/mitfusco98/HealthPrepV2/blob/ecb308aa61c685848a68ae5060ad34ca858be31c/services/jwt\\_service.py](https://github.com/mitfusco98/HealthPrepV2/blob/ecb308aa61c685848a68ae5060ad34ca858be31c/services/jwt_service.py)

55 66 **epic\_registration.html**

[https://github.com/mitfusco98/HealthPrepV2/blob/ecb308aa61c685848a68ae5060ad34ca858be31c/templates/admin/epic\\_registration.html](https://github.com/mitfusco98/HealthPrepV2/blob/ecb308aa61c685848a68ae5060ad34ca858be31c/templates/admin/epic_registration.html)

57 60 **migrate\_root\_admin.py**

[https://github.com/mitfusco98/HealthPrepV2/blob/ecb308aa61c685848a68ae5060ad34ca858be31c/migrate\\_root\\_admin.py](https://github.com/mitfusco98/HealthPrepV2/blob/ecb308aa61c685848a68ae5060ad34ca858be31c/migrate_root_admin.py)

71 72 74 **epic\_fhir\_service.py**

[https://github.com/mitfusco98/HealthPrepV2/blob/ecb308aa61c685848a68ae5060ad34ca858be31c/services/epic\\_fhir\\_service.py](https://github.com/mitfusco98/HealthPrepV2/blob/ecb308aa61c685848a68ae5060ad34ca858be31c/services/epic_fhir_service.py)

83 84 **keep\_alive.py**

[https://github.com/mitfusco98/HealthPrepV2/blob/ecb308aa61c685848a68ae5060ad34ca858be31c/keep\\_alive.py](https://github.com/mitfusco98/HealthPrepV2/blob/ecb308aa61c685848a68ae5060ad34ca858be31c/keep_alive.py)